

# MEMORY-AWARE SCHEDULING

JULIEN HERRMANN

ENSEEIH  
3SN-B - ALGORITHMES HPC

DECEMBER 7, 2023




- 1 Memory Usage and Performance on general task graphs
- 2 Tree-shaped task graphs
  - Sequential traversals
  - Parallel traversals
- 3 Backpropagation graphs
  - Single memory
  - Multiple memories
  - Parallel Processing

# **MEMORY USAGE AND PERFORMANCE ON GENERAL TASK GRAPHS**


- Decompose an application (simulations, scientific computations,...) into **tasks**
- Data produced and used by tasks create **dependencies**
- Task graph : Directed Acyclic Graph (DAG)
  - ▶ nodes : computational tasks
  - ▶ edges : data dependencies between tasks
- Task mapping and scheduling done at **runtime**
- Numerous runtime projects:
  - ▶ StarPU (Inria Bordeaux) : dynamically schedules tasks on any computing resource (CPU, GPU, \*PU)
  - ▶ DAGUE, ParSEC (ICL Tennessee) : task graph expressed in symbolic form for linear algebra
  - ▶ StarSs (Barcelona), Xkaapi (Grenoble), and others...

- Consider a simple task graph




# TASK GRAPH SCHEDULING AND MEMORY

- Consider a simple task graph
- Tasks have duration and memory demands




# TASK GRAPH SCHEDULING AND MEMORY

- Consider a simple task graph
- Tasks have duration and memory demands



# TASK GRAPH SCHEDULING AND MEMORY

- Consider a simple task graph
- Tasks have duration and memory demands




- **Peak Memory** : maximum memory usage



# TASK GRAPH SCHEDULING AND MEMORY


- Consider a simple task graph
- Tasks have duration and memory demands



- **Peak Memory** : maximum memory usage

# TASK GRAPH SCHEDULING AND MEMORY


- Consider a simple task graph
- Tasks have duration and memory demands



- **Peak Memory :** maximum memory usage

# TASK GRAPH SCHEDULING AND MEMORY

- Consider a simple task graph
- Tasks have duration and memory demands



- **Peak Memory** : maximum memory usage
- Trade-off between Peak Memory and Performance

# RESEARCH PROBLEMS

Several interesting questions:

- For **sequential processing** :
  - ▶ Minimum memory to process a graph
  - ▶ In case of memory shortage, minimum I/Os required
- In case of **parallel processing** :
  - ▶ Trade-off between memory and time
  - ▶ Makespan minimization under bounded memory


All of these problems are *NP-hard* on **general graphs**.

Sometimes restrict on simpler graphs:

- **Trees** (single output, multiple inputs for each task)  
Arise in sparse linear algebra (sparse direct solvers), with large data to handle : memory is a problem
- **Backpropagation graphs**  
Arise in automatic differentiation, gradient descent, training of deep neural networks,...


# SEQUENTIAL PROBLEM

- Tasks have no execution data
- Output data have a memory size





- Output data are kept in memory until every children is processed
- Even in the sequential case, scheduling influences the peak memory

## EXAMPLE



# EXAMPLE



Mem = 3

Exécution :



Exécution :

A

C

B


E

D

F

Mem Peak = 5

# EXAMPLE



Mem = 6

Exécution :





Exécution :

A C B E D F

Mem Peak = 5



# EXAMPLE



Mem = 5

Exécution :





Exécution :

A C B E D F

Mem Peak = 5

# EXAMPLE



Mem = 5

Exécution :





Exécution :

A C B E D F


Mem Peak = 5

# EXAMPLE



Mem = 4

Exécution :




Exécution :

A C B E D F

Mem Peak = 5

# EXAMPLE



Exécution :





Mem Peak = 6

Exécution :



A C B E D F

Mem Peak = 5


# EXAMPLE



# EXAMPLE



# EXAMPLE




Mem = 5

Exécution :




Mem Peak = 6

Exécution :



# EXAMPLE



Mem = 5

Exécution :





Mem Peak = 6

Exécution :






# EXAMPLE



# EXAMPLE



Exécution :  Mem Peak = 6

Exécution :  Mem Peak = 5

# RESEARCH PROBLEMS

Sequential processing of **general DAGs**:

- Finding the topological order that minimize peak memory on general DAGs is **NP-complete**
- The problem is still **NP-complete** on **unitary** edges (pebble game)


The problem becomes polynomial when we restrict on some simpler graphs:

- **Trees** (single output, multiple inputs for each task)  
Arise in sparse linear algebra (sparse direct solvers), with large data to handle : memory is a problem
- **Backpropagation graphs**  
Arise in automatic differentiation, gradient descent, training of deep neural networks,...

# TREE-SHAPED TASK GRAPHS


# MOTIVATION

In the **multifrontal method**, when factorizing a sparse matrix, the order in which variables can be eliminated is expressed with a bottom-up **elimination tree**.



- Any **topological order** of the elimination tree leads to a correct factorization
- Parallelism : separate subtrees can be processed in parallel
- Output data have **large size** (increasing closer to the root)

# NOTATIONS



- In-tree of  $n$  nodes
- Execution data of size  $m_i$
- Output data of size  $f_i$
- Leaf nodes have input data of null size

- Memory usage when executing node  $i$ :



$$MemReq(i) = \left( \sum_{j \in Children(i)} f_j \right) + m_i + f_i$$

# FROM IN-TREES TO OUT-TREES

## Theorem

*Considering an in-tree  $\mathcal{T}$  and a schedule  $\sigma_1$  of  $\mathcal{T}$ , we can build a schedule  $\sigma_2$  of the out-tree  $\overline{\mathcal{T}}$  obtained by reversing all edges, with the same peak memory:*

$$\sigma_2 = \text{reverse}(\sigma_1)$$




# POST-ORDER TRAVERSALS FOR TREE

**Post-order** : entirely process one subtree after the other  
(Deep First Search)

For each subtree  $\mathcal{T}_i$ :

- $P_i$  : peak memory
- $f_i$  : residual memory



For a given processing order of the subtrees  $\mathcal{T}_1, \mathcal{T}_2, \dots, \mathcal{T}_n$ ,  
the peak memory is:

$$\max \left\{ P_1; \right.$$




# POST-ORDER TRAVERSALS FOR TREE

**Post-order** : entirely process one subtree after the other  
(Deep First Search)

For each subtree  $\mathcal{T}_i$ :

- $P_i$  : peak memory
- $f_i$  : residual memory



For a given processing order of the subtrees  $\mathcal{T}_1, \mathcal{T}_2, \dots, \mathcal{T}_n$ ,  
the peak memory is:


$$\max \left\{ P_1; f_1 + P_2; \dots \right\}$$

# POST-ORDER TRAVERSALS FOR TREE

**Post-order** : entirely process one subtree after the other  
(Deep First Search)

For each subtree  $\mathcal{T}_i$ :

- $P_i$  : peak memory
- $f_i$  : residual memory



For a given processing order of the subtrees  $\mathcal{T}_1, \mathcal{T}_2, \dots, \mathcal{T}_n$ ,  
the peak memory is:


$$\max \left\{ P_1; f_1 + P_2; f_1 + f_2 + P_3; \dots; \right.$$

# POST-ORDER TRAVERSALS FOR TREE

**Post-order** : entirely process one subtree after the other  
(Deep First Search)

For each subtree  $\mathcal{T}_i$ :

- $P_i$  : peak memory
- $f_i$  : residual memory



For a given processing order of the subtrees  $\mathcal{T}_1, \mathcal{T}_2, \dots, \mathcal{T}_n$ ,  
the peak memory is:


$$\max \left\{ P_1; f_1 + P_2; f_1 + f_2 + P_3; \dots; \sum_{i < n} f_i + P_n; \right.$$

# POST-ORDER TRAVERSALS FOR TREE

**Post-order** : entirely process one subtree after the other  
(Deep First Search)

For each subtree  $\mathcal{T}_i$ :

- $P_i$  : peak memory
- $f_i$  : residual memory



For a given processing order of the subtrees  $\mathcal{T}_1, \mathcal{T}_2, \dots, \mathcal{T}_n$ ,  
the peak memory is:


$$\max \left\{ P_1; f_1 + P_2; f_1 + f_2 + P_3; \dots; \sum_{i < n} f_i + P_n; \sum_{i=1}^n f_i + m_R + f_R \right\}$$

# POST-ORDER TRAVERSALS FOR TREE

**Post-order** : entirely process one subtree after the other  
(Deep First Search)

For each subtree  $\mathcal{T}_i$ :

- $P_i$  : peak memory
- $f_i$  : residual memory



For a given processing order of the subtrees  $\mathcal{T}_1, \mathcal{T}_2, \dots, \mathcal{T}_n$ ,  
the peak memory is:

$$\max \left\{ \max_{j=1}^n \left( P_j + \sum_{i=1}^{j-1} f_i \right) ; \sum_{i=1}^n f_i + m_R + f_R \right\}$$

# LIU'S BEST POST-ORDER TRAVERSALS FOR TREES

## Theorem (Liu, Best Post-order Traversal)

*The best post-order traversal is obtain by processing subtrees in non-increasing order of  $P_i - f_i$ .*

## Proof by contradiction.

- Consider an optimal traversal which does not respect the order, that is to say:
  - ▶ subtree  $\mathcal{T}_j$  is processed right before subtree  $\mathcal{T}_k$
  - ▶ and  $P_k - f_k \geq P_j - f_j$
- Transform the schedule step by step without increasing the peak memory




# POST-ORDER IS NOT OPTIMAL

Theorem (Post-order Traversals are arbitrary bad in the general case)

*There is no constant  $K$  such that the best post-order traversal is a  $K$ -approximation in the general case.*

$\mathcal{T}_1$ :



- Minimum post-order peak memory:

$$BPO_1 = M + \epsilon + (b - 1)\frac{M}{b}$$


- Minimum traversal peak memory:

$$BT_1 = M + \epsilon + (b - 1)\epsilon$$

# POST-ORDER IS NOT OPTIMAL

Theorem (Post-order Traversals are arbitrary bad in the general case)

*There is no constant  $K$  such that the best post-order traversal is a  $K$ -approximation in the general case.*



- Minimum post-order peak memory:

$$BPO_2 = M + \epsilon + 2(b-1)\frac{M}{b}$$

- Minimum traversal peak memory:


$$BT_2 = M + \epsilon + 2(b-1)\epsilon$$



# POST-ORDER IS NOT OPTIMAL

Theorem (Post-order Traversals are arbitrary bad in the general case)

*There is no constant  $K$  such that the best post-order traversal is a  $K$ -approximation in the general case.*



- Minimum post-order peak memory:

$$BPO_K = M + \epsilon + \textcolor{red}{K}(b-1)\frac{M}{b}$$

- Minimum traversal peak memory:

$$BT_K = M + \epsilon + \textcolor{red}{K}(b-1)\epsilon$$

- Thus:

$$\frac{BPO_K}{BT_K} > K$$

# POST-ORDER IS NOT OPTIMAL

Theorem (Post-order Traversals are arbitrary bad in the general case)


*There is no constant  $K$  such that the best post-order traversal is a  $K$ -approximation in the general case.*

The best post-order is not optimal in the general case but **efficient** in practice:

	actual assembly trees	random trees
Non-optimal traversals	4.2%	61%
Maximum increase compared to optimal	18%	22%
Average increase compared to optimal	1%	12%

# MODEL FOR PARALLEL TREE PROCESSING

- $P$  uniform processors
- Shared memory of size  $M$
- Task  $i$  has execution times  $w_i$
- Simultaneous processing of nodes induces larger memory
- Trade-off time vs. memory



When processing a tree with multiple processors, there are multiple sources of parallelism:

- **Node parallelism** : a task node can be processed using multiple processors.  
⇒ It does not increase the memory usage but induces a lot of communications between processors.
- **Tree parallelism** : independent tasks can be processed at the same time by different processors  
⇒ It increases the memory usage since their data coexists at the same time in the shared memory.

# COMPLEXITY RESULTS

For **tree-shaped** task graph:

- Makespan minimization is NP-complete for general trees
- Makespan minimization is polynomial for unit-weight task trees
- Memory minimization is polynomial for  $w_i = 1$ ,  $m_i = 0$ , and  $f_i = 1$  (pebble game model)

## Theorem


*Deciding whether a tree can be scheduled with a bound  $M$  on memory and a bound  $C$  on makespan is NP-complete*

## Theorem

*There is no algorithm that is both an  $\alpha$ -approximation for makespan minimization and a  $\beta$ -approximation for peak memory minimization when scheduling tree-shaped task graphs*

# ALL-TO-ALL MAPPING


**All-to-all mapping:** post-order traversal of the tree, where all the processors work at every node (maximum node parallelism on every node)



- $T_1, T_2, \dots, T_n$  are processed in the best post-order using a all-to-all mapping scheduling
- Every processor executes root  $R$  in parallel
- Optimal memory scalability : same peak memory as the sequential execution
- No tree parallelism
- A lot of communications between processors

# PROPORTIONAL MAPPING

**Proportional mapping:** every subtrees are processed in parallel by a subset of processors proportional to their work load




- Subtree  $\mathcal{T}_i$  is executed by  $p_i$  processors where:

$$p_i = P * \frac{W_{\mathcal{T}_i}}{W_{\mathcal{T}}} \text{ and}$$
$$W_{\mathcal{T}} = \sum_{node \in \mathcal{T}} W_{node}$$

- Good work-load balance to exploit tree parallelism
- Memory scalability can be arbitrary bad compared to the sequential execution:  $P * M_{max}(P) \gg M_{seq}$

# MEMORY-AWARE MAPPINGS

**Memory-aware mapping:** (Agullo et al. [3]): aims at enforcing a given memory bound  $M_B$  on the peak memory



- Try to apply proportional mapping
- Check whether enough memory for each tree. If not, serialize them and update  $M_B$ :

- Ensures the given memory constraint and provides reliable estimates
- Tends to assign many processors on nodes at the top of the tree  $\Rightarrow$  performance issues on parallel nodes.



For **parallel** traversals of **tree-shaped** tasks graphs:

- Optimizing both memory and makespan is NP-complete
- Optimizing makespan under memory constraint is NP-complete
- No scheduling algorithm can be a constant factor approximation on both memory and time

Use of heuristics:

- All-to-all mapping : full node parallelism, no tree parallelism
- Proportional mapping : tree parallelism needing more memory
- Other memory-aware mappings

# BACKPROPAGATION GRAPHS

# AUTOMATIC DIFFERENTIATION

## Ice-sheet model:

### Model Algorithm (single timestep)

1. Evaluate driving stress  $\tau_d = \rho g h \nabla s$
2. Solve for velocities  
DO  $i = 1, \text{max\_iter}$ 
  - i. Evaluate nonlinear viscosity  $v_i$  from iterate  $u_i$
  - ii. Construct stress matrix  $A\{v_i\}$
  - iii. Solve linear system  $A u_{i+1} = \tau_d$
  - iv. (Exit if converged)ENDDO
3. Evolve thickness (continuity eqn)  
Automatic differentiation (AD) tools generate code for adjoint of operations

## Simpler Version:

```
proc Model Algorithm( $x_0$ )  
begin  
  Do stuff;  
  for  $i = 0$  to  $n$  do  
     $x_{i+1} = f_i(x_i)$ ;  
    Do stuff;  
  end  
  \(*F(u_0) = f_n \circ f_{n-1} \circ \dots \circ f_0(u_0)*\  
  Compute  $\nabla F(x_0) \cdot y$ ;  
end
```

## A quick reminder about the gradient:


$$\begin{aligned} F(u_0) &= f_n \circ f_{n-1} \circ \dots \circ f_1 \circ f_0(u_0) \\ \nabla F(u_0) \mathbf{y} &= Jf_0(u_0)^T \cdot \nabla(f_n \circ f_1)(u_1) \cdot \mathbf{y} \\ &= Jf_0(u_0)^T \cdot Jf_1(u_1)^T \cdot \dots \cdot Jf_{n-1}(u_{n-1})^T \cdot Jf_n(u_n)^T \cdot \mathbf{y} \end{aligned}$$

$$\begin{aligned} Jf^T &= \text{Transpose Jacobian matrix of } f; \\ u_{i+1} &= f_i(u_i) = f_i(f_{i-1} \circ \dots \circ f_0(u_0)). \end{aligned}$$

# TASK GRAPH DESCRIPTION


$$\mathbf{F}_i(\mathbf{x}_i) = \mathbf{x}_{i+1} \quad i < l \quad \text{(Forward Phase)}$$

$$B_i(x_i, y_{i+1}) = y_i \quad i \leq l \quad \text{(Backward Phase)}$$



# TASK GRAPH DESCRIPTION


$$F_i(x_i) = x_{i+1} \quad i < l \quad (\text{Forward Phase})$$
$$B_i(x_i, y_{i+1}) = y_i \quad i \leq l \quad (\text{Backward Phase})$$



# TASK GRAPH DESCRIPTION

$$F_i(x_i) = x_{i+1} \quad i < l \quad (\text{Forward Phase})$$


$$B_i(x_i, y_{i+1}) = y_i \quad i \leq l \quad (\text{Backward Phase})$$



# RELATION TO DEEP LEARNING

When training a neural network:

- **Forward phase:** computes predicted output for each layers with respect to model weights
- **Loss computation:** difference between predicted output and expected output
- **Backward phase:** updates the model weights to minimize loss function



GoogleNet graph

# RELATION TO DEEP LEARNING


When training a neural network:

- **Forward phase:** computes predicted output for each layers with respect to model weights
- **Loss computation:** difference between predicted output and expected output
- **Backward phase:** updates the model weights to minimize loss function






# BACKPROPAGATION GRAPHS




- No graph parallelism (linear structure)
- Intermediate data ( $x_i$  and  $y_i$ ) have large sizes
- Intermediate data can be useful much later in execution
- Intermediate data **can not all fit** in memory at the same time
- Initial state :  $x_0$  is stored in memory
- **Objective** : compute  $y_0$


# MODEL OF EXECUTION : STORE ALL



Memory :




# MODEL OF EXECUTION : STORE ALL



Memory :

$x_0$	$x_1$				
-------	-------	--	--	--	--



# MODEL OF EXECUTION : STORE ALL




Memory :

$x_0$	$x_1$	$x_2$			
-------	-------	-------	--	--	--



# MODEL OF EXECUTION : STORE ALL




# MODEL OF EXECUTION : STORE ALL



# MODEL OF EXECUTION : STORE ALL



# MODEL OF EXECUTION : STORE ALL




Memory : 

$x_0$	$x_1$	$x_2$	$x_3$	$x_4$	$y_5$
-------	-------	-------	-------	-------	-------




# MODEL OF EXECUTION : STORE ALL



Memory : 

$x_0$	$x_1$	$x_2$	$x_3$	$y_4$	
-------	-------	-------	-------	-------	--


# MODEL OF EXECUTION : STORE ALL





Memory :

$x_0$	$x_1$	$x_2$	$y_3$		
-------	-------	-------	-------	--	--



# MODEL OF EXECUTION : STORE ALL



# MODEL OF EXECUTION : STORE ALL



# MODEL OF EXECUTION : STORE ALL




For  $l = 5$  forward steps

Strategy **Store all** (memory expensive):

- Peak Memory : 6
- Recomputation : 0


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory :

$x_0$		
-------	--	--

# MODEL OF EXECUTION : RECOMPUTE ALL




Memory :

$x_0$	$x_1$	
-------	-------	--




# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_2$	
-------	-------	--


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory :

$x_0$	$x_3$	
-------	-------	--


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_4$	
-------	-------	--


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_5$	
-------	-------	--


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$		$y_5$
-------	--	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory :

$x_0$	$x_1$	$y_5$
-------	-------	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_2$	$y_5$
-------	-------	-------

# MODEL OF EXECUTION : RECOMPUTE ALL




Memory :

$x_0$	$x_3$	$y_5$
-------	-------	-------




# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_4$	$y_5$
-------	-------	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$		$y_4$
-------	--	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_1$	$y_4$
-------	-------	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_2$	$y_4$
-------	-------	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_3$	$y_4$
-------	-------	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$		$y_3$
-------	--	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory :

$x_0$	$x_1$	$y_3$
-------	-------	-------

# MODEL OF EXECUTION : RECOMPUTE ALL




Memory :

$x_0$	$x_2$	$y_3$
-------	-------	-------




# MODEL OF EXECUTION : RECOMPUTE ALL



Memory :

$x_0$		$y_2$
-------	--	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$x_0$	$x_1$	$y_2$
-------	-------	-------


# MODEL OF EXECUTION : RECOMPUTE ALL



Memory :

$x_0$		$y_1$
-------	--	-------

# MODEL OF EXECUTION : RECOMPUTE ALL



Memory : 

$y_0$		
-------	--	--

For  $l = 5$  forward steps


Strategy **Store all** (memory expensive):

- Peak Memory : 6
- Recomputation : 0

Strategy **Recompute all** (compute expensive):

- Peak Memory : 3
- Recomputation : 10


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory :

$x_0$			
-------	--	--	--


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory :

$x_0$	$x_1$		
-------	-------	--	--

## EXECUTION : STORE SOME / RECOMPUTE SOME




Memory :

$x_0$	$x_2$		
-------	-------	--	--




## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$	$x_3$	
-------	-------	-------	--


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$	$x_4$	
-------	-------	-------	--


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$	$x_5$	
-------	-------	-------	--


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory :

$x_0$	$x_2$		$y_5$
-------	-------	--	-------


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$	$x_3$	$y_5$
-------	-------	-------	-------


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$	$x_4$	$y_5$
-------	-------	-------	-------


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$		$y_4$
-------	-------	--	-------

## EXECUTION : STORE SOME / RECOMPUTE SOME




Memory : 

$x_0$	$x_2$	$x_3$	$y_4$
-------	-------	-------	-------




## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$	$x_2$		$y_3$
-------	-------	--	-------


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory : 

$x_0$			$y_2$
-------	--	--	-------


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory :

$x_0$	$x_1$		$y_2$
-------	-------	--	-------


## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory :

$x_0$			$y_1$
-------	--	--	-------

## EXECUTION : STORE SOME / RECOMPUTE SOME



Memory :

$y_0$			
-------	--	--	--

For  $l = 5$  forward steps

Strategy **Store all** (memory expensive):

- Peak Memory : 6
- Recomputation : 0

Strategy **Recompute all** (compute expensive):

- Peak Memory: 3
- Recomputation : 10

Strategy **Store some / Recompute some** (hybrid):

- Peak Memory: 4
- Recomputation : 4

# SCHEDULING PROBLEM

## Application Parameters:

- $l$  : number of forward steps in the BP graph
- $x_i$  : memory size of intermediate value  $i$
- $wf, wb$  : computational cost of forward and backward steps

## Memory Parameters:

- $w_m$  : writing cost in memory
- $r_m$  : reading cost in memory
- $c_m$  : size of the memory

**Question :** Which intermediate data should we store in memory and which should we evict and recompute later?

# HOMOGENEOUS MODEL + SINGLE FREE MEMORY

## Application Parameters:

- $l$  : arbitrary number of steps
- $\mathbf{x}_i = \mathbf{1}$  : unitary size of intermediate value
- $\mathbf{w}\mathbf{f}_i = \mathbf{1}, \mathbf{w}\mathbf{b}_i = \mathbf{1}$  : homogeneous forward and backward steps

## Memory Parameters:

- $\mathbf{w}_m = \mathbf{0}$  : free writing cost in memory
- $\mathbf{r}_m = \mathbf{0}$  : free reading cost in memory
- $c_m$  : arbitrary memory size

**Griewank and Walther, 2000:** *Revolve*( $l, c_m$ ), optimal algorithm with  $c_m$  memory slots on homogeneous backpropagation graphs and free memory



# HETEROGENEOUS MODEL + $K$ MEMORIES

We consider  $K$  different memories with arbitrary reading and writing costs:

## Application Parameters:

- $l$  : arbitrary number of steps
- $x_i$  : arbitrary memory size of intermediate value
- $wf_i, wb_i$  : arbitrary computational cost for forward and backward steps

## Memory Parameters:

- $w_m^{(k)}$  : writing cost into memory  $k$
- $r_m^{(k)}$  : reading cost from memory  $k$
- $c_m^{(k)}$  : size of memory  $k$

**Aupy, Herrmann, Hovland, Robert, 2015:** Optimal algorithm for two level of storage: cheap bounded memory and costly unbounded disks.

**Aupy, Herrmann, 2019:** Library of optimal schedules for any number of storage level.

# PARALLEL PROCESSING OF BACKPROPAGATION GRAPHS

There are multiple parallelization techniques for Deep Neural Networks:

- **Model parallelism:** the model is partitioned on the architecture
  - ▶ Intra-layer parallelism : partition individual layers across workers
  - ▶ Inter-layer parallelism : pipelining
  - ▶ ...
- **Data parallelism:** the model is replicated on several workers, and each worker computes a micro-batch
  - ▶ ZeroDP
  - ▶ Fully-shared DP
  - ▶ ...

How do we optimize memory usage in these parallel frameworks to train deeper networks or bigger batches? ⇒ **INTERNSHIP**

# REFERENCES

[1] Liu, J. W.

**On the storage requirement in the out-of-core multifrontal method for sparse factorization.**

*ACM Transactions on Mathematical Software (TOMS)*, 1986

[2] Liu, J. W.

**An application of generalized tree pebbling to sparse matrix factorization.**

*SIAM Journal on Algebraic Discrete Methods*, 1987

[3] Agullo, E., Guermouche, A., L'Excellent, J. Y.

**Reducing the I/O volume in sparse out-of-core multifrontal methods.**

*SIAM Journal on Scientific Computing*, 2010

[4] Griewank, A., Walther, A.

**Revolve: an implementation of checkpointing for the reverse or adjoint mode of computational differentiation.**

*ACM Transactions on Mathematical Software (TOMS)*, 2000

[5] Aupy, G., Herrmann, J., Hovland, P., Robert, Y.

**Optimal multistage algorithm for adjoint computation.**

*SIAM Journal on Scientific Computing*, 2016

[6] Aupy G., Herrmann, J.

**H-Revolve: a framework for adjoint computation on synchronous hierarchical platforms.**

*ACM Transactions on Mathematical Software (TOMS)*, 2020

[7] Huang, Y., Cheng, Y., Bapna, A., Firat, O., Chen, D., Chen, M., ...

**Gpipe: Efficient training of giant neural networks using pipeline parallelism.**

*Advances in neural information processing systems*, 2019